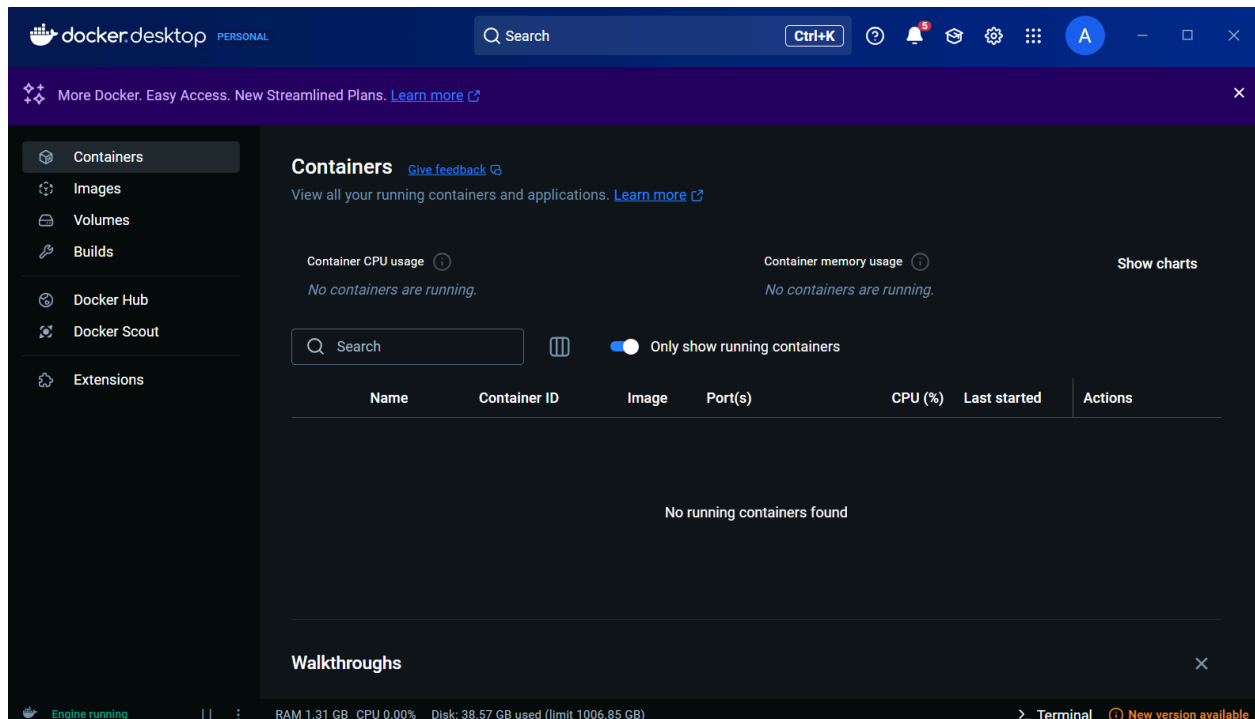


# ENVIRONMENTAL SETUP AND LEARNING PROCESS

## DOCKER AND ENVIRONMENT SETUP

To establish a robust development environment with all dependencies and ensure portability, Docker is used. The goal was to create an environment where any change in the codebase reflects in the project without manual intervention, as this is a full-fledged product rather than a single-file program. Docker Compose plays a crucial role in achieving this by managing multiple containers efficiently.

Docker Desktop is installed from the website. Then, the project folders are cloned from the GitHub repository, and Docker Compose is used to set up the environment, ensuring that pgwatch, Grafana, and the PostgreSQL database are running for monitoring.



```
git clone https://github.com/cybertec-postgresql/pgwatch.git && cd pgwatch
docker compose -f ./docker/docker-compose.yml up --detach
```

To gain a deeper understanding of the PostgreSQL server, the pgAdmin container is also included:

**docker compose -f ./docker/docker-compose.yml up --detach pgadmin**

The different components of the project run on various ports:

PgAdmin: port 80

Pgwatch: port 8080

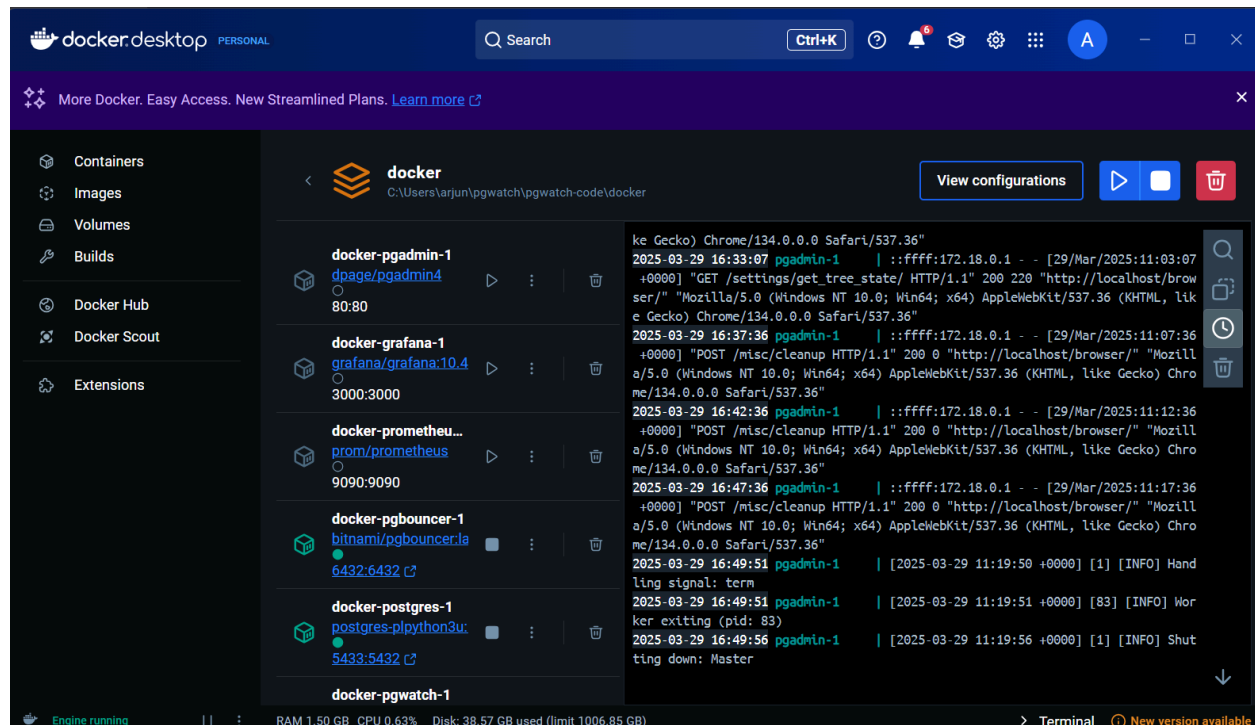
Grafana: port 3000

Prometheus: port 9090

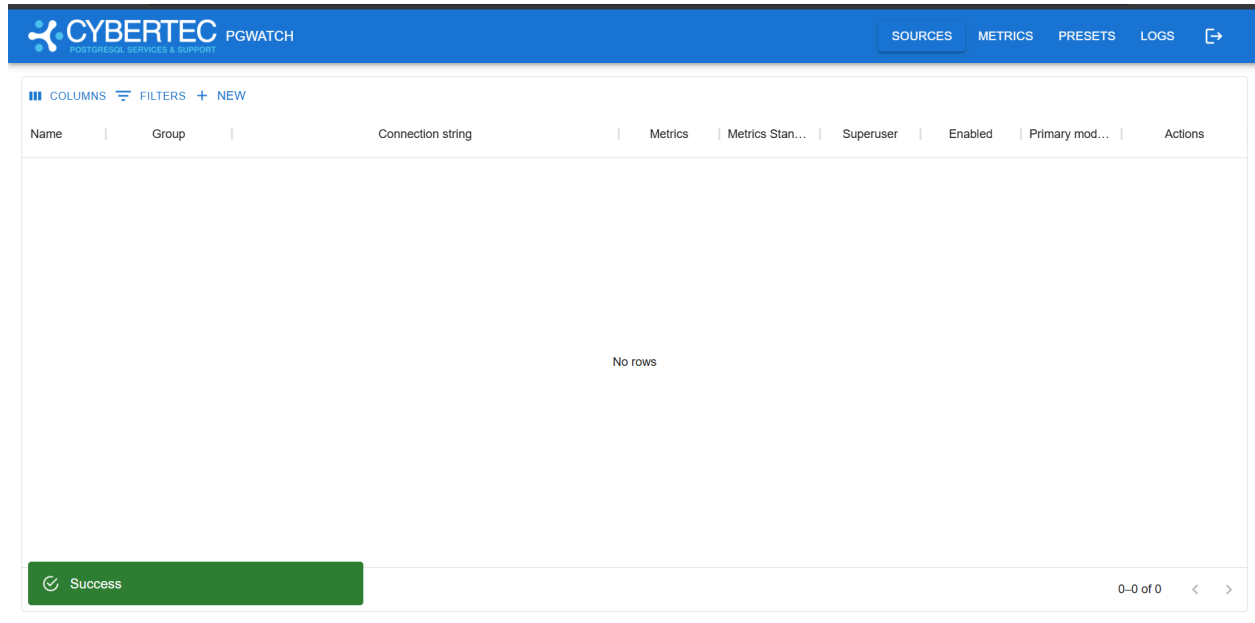
PostgreSQL: port 5433 (originally on 5432, but changed locally to 5433 to accommodate pgwatch)

Using:

**docker compose -f ./docker/docker-compose.yml up --detach**



we effectively set up a development environment where changes in the codebase are reflected in the running application without the need for manual reconfiguration, streamlining the development process.

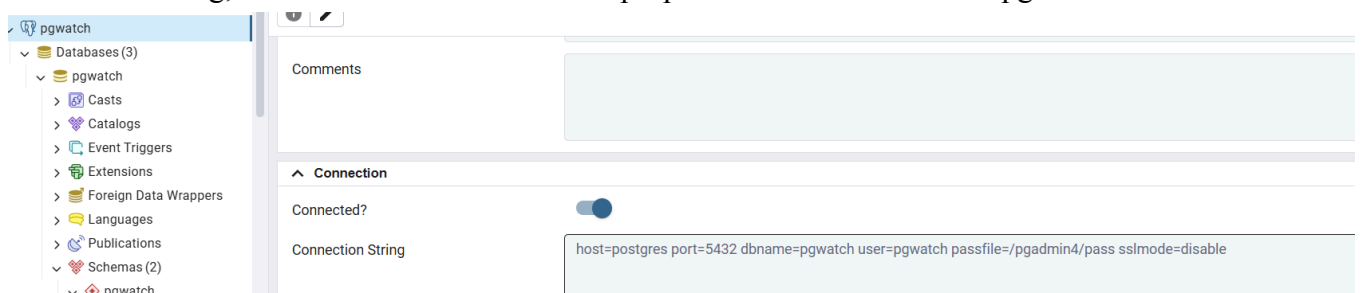


The pgwatch container is changed to run 5433 on the local ports by modifying the compose.postgres.yml file in the docker folder.

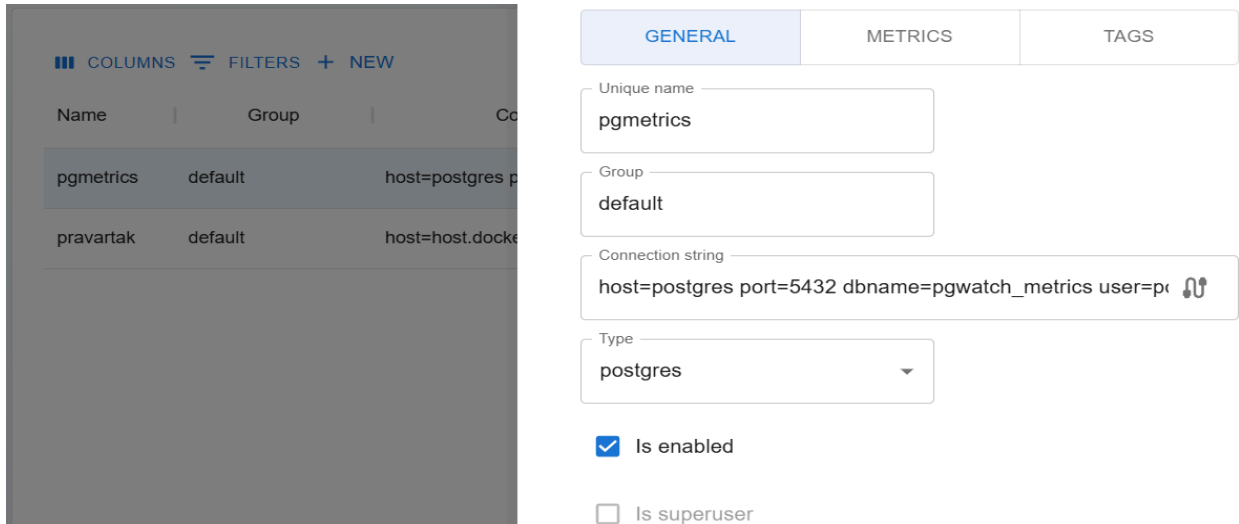
```
C: > Users > arjun > pgwatch > docker > compose.postgres.yml
1  services:
2    postgres:
12   command:
15     - "-ctrack_io_timing=on"
16     - "-ctrack_functions=pl"
17   ports:
18     - "5433:5432"
```

## CONNECTING DATABASES TO BE MONITORED

Two databases were connected with pgwatch to be monitored. We can connect them using a connection string, which was obtained from the properties of the databases in pgadmin.



This connection string is given in pgwatch(ui) to connect the database to be monitored. This is done for a database in the postgres container we installed.



The screenshot shows the pgwatch web interface. On the left, a table lists monitored databases:

Name	Group	Connection string
pgmetrics	default	host=postgres p...
pravartak	default	host=host.docke...

On the right, the configuration form for the 'pgmetrics' entry is shown under the 'GENERAL' tab:

- Unique name: pgmetrics
- Group: default
- Connection string: host=postgres port=5432 dbname=pgwatch\_metrics user=pgwatch
- Type: postgres
- Is enabled: ☒
- Is superuser: ☐

To connect a database from the postgres locally installed, which is another port, through research the below given modification to hostname was used.



The screenshot shows the configuration form for the 'pravartak' entry. The connection string is modified to:

host=host.docker.internal port=5432 dbname=pravartak user=pgwatch

As a base requirement the following instructions are run on the database for connecting and fetching metrics.

```
CREATE ROLE pgwatch WITH LOGIN PASSWORD 'secret';  
-- For critical databases it might make sense to ensure that the user account  
-- used for monitoring can only open a limited number of connections  
-- (there are according checks in code, but multiple instances might be launched)  
ALTER ROLE pgwatch CONNECTION LIMIT 5;  
GRANT pg_monitor TO pgwatch;  
GRANT CONNECT ON DATABASE mydb TO pgwatch;  
GRANT EXECUTE ON FUNCTION pg_stat_file(text) to pgwatch; -- for wal_size metric  
GRANT EXECUTE ON FUNCTION pg_stat_file(text, boolean) TO pgwatch;
```



Here we can see how changing the database(pgwatch) can be seen in the ui of pgwatch.

```
pgwatch=> update metric set description='updated' where name='archiver';
UPDATE 1
```

Name	Description	InitSQL	Node status	Gauges	Instance level?	Storage name
archiver	updated			is_falling_int,seconds_sin...	✓	
backends				*		
backup_age_pgb...		CREATE EXTENSION IF ...			✓	
backup_age_walg		CREATE EXTENSION IF ...			✓	
bgwriter			primary		✓	

The database pgwatch\_metrics contain the the tables where the recorded metrics are stored, like db\_stats,db\_size,stat\_statements\_calls etc where different live metrics that are recorded are stored.

```
You are now connected to database "pgwatch_metrics" as user "pgwatch".
pgwatch_metrics=> \dt
```

List of relations				
Schema	Name	Type	Owner	
public	archiver	partitioned table	pgwatch	
public	backends	partitioned table	pgwatch	
public	bgwriter	partitioned table	pgwatch	
public	change_events	partitioned table	pgwatch	
public	checkpointer	partitioned table	pgwatch	
public	configuration_changes	partitioned table	pgwatch	
public	configured_dbs	partitioned table	pgwatch	
public	cpu_load	partitioned table	pgwatch	
public	db_size	partitioned table	pgwatch	
public	db_stats	partitioned table	pgwatch	
public	index_changes	partitioned table	pgwatch	

Each of these contain four columns, time,dbname,data and tagdata. Thus we can see that it is a time series database that stores the recorded data from each table at intervals of time.

Below given are some examples of recorded data

db\_stats:

	time timestamp with time zone	dbname text	data jsonb
1	2025-03-21 12:45:04.328285+00	pgmetrics	{"sys_id": "7483764217247694875", "blks_hit": 30290, "sessions": 27, "blks_rea
2	2025-03-21 12:46:09.807231+00	pgmetrics	{"sys_id": "7483764217247694875", "blks_hit": 38692, "sessions": 28, "blks_rea
3	2025-03-21 12:47:15.363337+00	pgmetrics	{"sys_id": "7483764217247694875", "blks_hit": 42009, "sessions": 31, "blks_rea
4	2025-03-21 12:48:20.853984+00	pgmetrics	{"sys_id": "7483764217247694875", "blks_hit": 44488, "sessions": 34, "blks_rea
5	2025-03-21 12:49:26.327884+00	pgmetrics	{"sys_id": "7483764217247694875", "blks_hit": 47013, "sessions": 37, "blks_rea
6	2025-03-21 12:50:31.717514+00	pgmetrics	{"sys_id": "7483764217247694875", "blks_hit": 51033, "sessions": 39, "blks_rea

stat\_statements\_calls:

	time timestamp with time zone	dbname text	data jsonb
1	2025-03-28 07:28:54.073532+00	pgmetrics	{"calls": 1643, "sys_id": "7483764217247694875", "total_time": 969.332, "real_d
2	2025-03-28 07:29:59.87421+00	pgmetrics	{"calls": 2133, "sys_id": "7483764217247694875", "total_time": 1243.451, "real_
3	2025-03-28 07:31:05.846601+00	pgmetrics	{"calls": 2235, "sys_id": "7483764217247694875", "total_time": 1253.156, "real_
4	2025-03-28 07:32:11.697991+00	pgmetrics	{"calls": 2352, "sys_id": "7483764217247694875", "total_time": 1297.566, "real_
5	2025-03-28 07:33:17.628176+00	pgmetrics	{"calls": 2503, "sys_id": "7483764217247694875", "total_time": 1342.008, "real_
6	2025-03-28 07:34:23.634523+00	pgmetrics	{"calls": 2631, "sys_id": "7483764217247694875", "total_time": 1366.202, "real_

From exploring, it seems as if grafana shows the metrics like tps, qps etc by querying this databases with the appropriate functions to get the values for them.

## BASIC UNDERSTANDING OF SOURCE CODE

### Frontend (Web UI)

- Framework: React.js (v18.2.0)
- Language: TypeScript
- UI Components: Material-UI (MUI) v5

### Backend

Language: Go (Golang)

Architecture: Microservices-based with several components:

- webserver: Main API server
- metrics: Metrics collection and processing
- sources: Data sources
- sinks: Data output handlers
- reaper: Cleanup/maintenance tasks
- db: Database interactions